

Content Understanding and Document Workflow

Customer discussion deck

Turn document understanding into a governed business workflow.

Use this deck to agree on the decision, automation evidence, and controls needed before deployment.

How to use this conversation

Use this working-module deck with sponsors, SMEs, security, data owners, and engineering. It is not an implementation walkthrough.

Each module moves through the same three steps:

- Discuss: Why the decision matters to the business outcome.
- Decide: Which path fits the documents, workflow, and operating environment.
- Prove: What the team must show before moving forward.

Record the choice, owner, open question, and evidence gate as you go. The practical steps live in the scenario modules.

Module 1: Confirm scope and connect the foundation

Agree the document class and downstream action. Select the intended extraction path before provisioning, and reuse approved tenant resources.

Discuss:

- Which business decision this workflow will improve.
- Which teams own identity, networking, storage, monitoring, and model access.
- Which environments are needed for experimentation, pilots, and production.
- How to apply keyless access, least privilege, and traceability from the start.

Build a reusable base that later scenario tracks can share. Avoid a one-off demo environment.

Foundation choices and trade-offs

Key decisions:

- Shared vs. dedicated resources: Shared foundations reduce setup work; dedicated resources can simplify isolation and chargeback.
- Region and model availability: Consider model choice, data residency, latency, and quota together.
- Identity model: Prefer managed identity and role-based access over keys where possible.
- Observability baseline: Start tracing and monitoring early so later workflow issues are diagnosable.

The trade-off is simple: quick setup must not hide production controls needed later.

Foundation evidence

Before the team proceeds, confirm:

- The Foundry foundation, storage, model deployments, and monitoring plan have owners.
- The access model is documented and keyless-first.
- Environment boundaries are understood.
- Later modules can use required outputs without copying secrets into notes or code.
- The team knows what still needs security, networking, or operations review.

Decision question: Can engineering safely build on this foundation without re-deciding basic platform controls each module?

Module 2: Connect an approved document source

Content Understanding starts with approved documents, not files that merely happen to be available.

Discuss:

- Which workflow source is authoritative: Azure Blob, ADLS Gen2, SharePoint, OneLake, or another governed store.
- Who is allowed to approve sample use.
- How source identity, document version, permissions, and retention are preserved.
- How unsafe, unsupported, or out-of-scope files are quarantined.

A clean folder of copied files works for a lab. It is not an approved business source.

Document source choices and trade-offs

Key decisions:

- Business source vs. staging area: Direct integration preserves context; staging can simplify processing but adds governance work.
- Representative samples: Include normal, edge, low-quality, multilingual, and failure cases, not only ideal examples.
- Permissions: Decide whether the workflow inherits source permissions or uses a separate processing identity.
- Retention: Define how originals, extracted results, evidence, and corrections are retained.

Faster intake does not justify losing provenance or authorization.

Document source evidence

Before moving forward, confirm:

- The document source is approved for this scenario.
- Sample documents are representative and authorized for testing.
- Each document traces to its source, version, owner, and retention policy.
- The team has a clear path for quarantine, deletion, and exception handling.
- Access boundaries are enforceable, not just assumed.

Decision question: Can every document used by the workflow be explained, traced, and governed?

Module 3: Select the extraction capability

Choose extraction based on document variability, schema needs, confidence requirements, and operating constraints.

Discuss:

- Whether Content Understanding analyzers fit the document classes and fields.
- Where Document Intelligence, LLM structured outputs, or multimodal approaches may fit better.
- Which fields require exact evidence versus broad summarization.
- Which errors are tolerable, reviewable, or unacceptable.

The customer should leave with a documented capability decision, not a default product choice.

Extraction choices and trade-offs

Key decisions:

- Content Understanding: Fits SMEs who need to shape classes and schemas for varied content.
- Document Intelligence: Useful for established document extraction patterns and form-like structure.
- LLM structured outputs: Flexible for reasoning over text but require strict validation and evidence controls.
- Multimodal processing: Helpful when layout, images, or visual cues matter.

More flexible extraction increases validation, review, cost, and monitoring work.

Extraction selection evidence

Before implementation, confirm:

- The selected capability matches document quality, field complexity, region, cost, and review needs.
- The target schema is specific enough to test.
- Known failure cases are included in the decision.
- The team has agreed when a field must be empty rather than inferred.
- Human-review rules are part of the extraction decision.

Decision question: Can the team explain why this capability is the right fit for the first controlled workflow?

Module 4: Implement typed extraction with evidence

A useful extraction result is typed, validated, and supported by evidence.

Discuss:

- Which fields are required, optional, derived, or prohibited.
- What evidence is needed for each important value: page, span, citation, confidence, or source reference.
- How missing, ambiguous, conflicting, or low-confidence values should be represented.
- Where the workflow must avoid inferred values.

Build more than a JSON shape. Build a decision record a reviewer and auditor can trust.

Typed extraction choices and trade-offs

Key decisions:

- Strict schema vs. flexible notes: Strict schemas help automation; flexible notes may help SMEs explain unusual cases.
- Confidence thresholds: High thresholds reduce false approvals and increase review volume.
- Evidence granularity: More detailed evidence improves trust but can increase storage and UI complexity.
- Failure behavior: Empty-with-reason is safer than filling a field without support.

Automation depends on handling uncertainty well, not only on high field coverage.

Typed extraction evidence

Before the workflow can use extracted results, confirm:

- The output validates against the agreed schema.
- Important fields include evidence and confidence where appropriate.
- Missing and low-confidence fields follow a consistent policy.
- Unsupported values are not invented to satisfy the schema.
- Reviewers can see enough source context to confirm or challenge a value.

Decision question: Would a business reviewer understand what was extracted, why it was trusted, and what still needs attention?

Module 5: Build review, correction, and handoff

Human review belongs in the product design. It cannot be a fallback after automation fails.

Discuss:

- Which cases require review: missing fields, low confidence, conflicting values, sensitive decisions, or policy exceptions.
- Who can approve, correct, reject, or escalate.
- What correction reason and evidence must be retained.
- Which downstream system receives approved results.

The review experience should make the right action easier than bypassing the process.

Review and handoff choices and trade-offs

Key decisions:

- Reviewer queue vs. embedded workflow: Queues centralize review; embedded workflows meet users where they already work.
- Correction model: Corrections should update the case record and feed evaluation. Do not silently overwrite history.
- Approval boundary: Decide which results can flow automatically and which require a named approver.
- Handoff seam: Use a governed downstream contract. Do not write directly from unreviewed extraction.

Reducing reviewer effort must not erase accountability.

Review evidence

Before handoff is trusted, confirm:

- Review rules are explicit and testable.
- A reviewer can correct, approve, reject, and explain the decision.
- Corrections retain the document, evidence, reviewer, timestamp, and contract version.
- Approved results are handed off through a controlled interface.
- Exceptions have an owner and a resolution path.

Decision question: Can the customer prove who approved a result, what changed, and why it was sent downstream?

Module 6: Evaluate and trace the workflow

A document workflow is ready when it performs consistently on representative cases and failures are diagnosable.

Discuss:

- Which quality measures matter: field accuracy, routing accuracy, false approvals, review rate, latency, and cost.
- Which adversarial or messy cases should be tested.
- What traces must show across intake, extraction, review, and handoff.
- How corrections become future evaluation evidence.

Evaluate real workflow risk, not only model output quality.

Evaluation and tracing choices and trade-offs

Key decisions:

- Holdout data: Keep evaluation examples separate from tuning examples.
- Quality gates: Define thresholds for automation, review, and rejection.
- Trace detail: Capture enough context to debug without exposing unnecessary sensitive content.
- Regression testing: Re-run important cases when schemas, analyzers, prompts, or review rules change.

More automation without measurement increases hidden business risk.

Evaluation evidence

Before promotion, confirm:

- Representative test cases cover normal, edge, and failure segments.
- Field accuracy, review volume, false approval risk, and latency are measured.
- Traces connect document intake, extraction, policy, review, correction, and handoff.
- Failures have named causes and owners.
- The release decision covers rollback and monitoring.

Decision question: Can the team defend the workflow with evidence rather than a successful demo?

Module 7: Deploy the reviewable workflow

Deployment turns an experiment into an accountable business service.

Discuss:

- Who can invoke the workflow and under what identity.
- Which environment receives the pilot and what production controls are required.
- How analyzer or schema versions are promoted.
- How monitoring, support, rollback, and change management work.

Keep the first deployment controlled, observable, and reversible.

Deployment choices and trade-offs

Key decisions:

- Pilot scope: Start with a bounded document set, user group, and decision path.
- Endpoint and identity: Use authenticated access and managed identity where possible.
- Versioning: Track analyzer, schema, policy, review rules, and downstream contract together.
- Operations: Decide alert ownership, support process, rollback criteria, and release cadence.

A broad rollout before operational readiness can create more manual work than it removes.

Deployment evidence

Before controlled rollout, confirm:

- The workflow runs behind approved access controls.
- Monitoring and tracing are enabled for the end-to-end path.
- Version and rollback information is documented.
- Review and correction data stay available after deployment.
- The pilot owner can decide whether to expand, pause, or revise the workflow.

Decision question: Is the workflow ready to serve a bounded real use case with accountable controls?

Close the discussion

Before the next working session, record:

- the first high-value document decision and its business owner
- the approved source and 15–30 safe, representative samples
- expected fields, unacceptable errors, and review rules
- the SME, engineering, security, and workflow handoff owners
- current constraints for identity, retention, monitoring, and deployment
- one owner and due date for every unresolved decision

Agree on the first pilot slice and the evidence needed before expansion.