

Avatar Scenario

An application with accessible output

Build an interactive assistant or a repeatable content-production application in the customer's tenant. Start with the application goal, then choose whether an avatar helps.

Choose the application goal, media path, approval boundary, accessible fallback, and evidence needed for a controlled pilot.

How to use this conversation

Each module moves through the same three steps:

Step	Customer question
Discuss	Which user or content outcome are we trying to improve?
Decide	Which experience and control path fits the risk?
Prove	What must the customer see before the team moves on?

Record the choice, owner, open question, and evidence gate as you go. Keep build detail in the scenario modules.

Module 1. Context

Choose the application goal

Interactive assistant: help users ask questions or complete a task using approved knowledge.

Content-production application: turn approved source revisions into media through a repeatable generation and publishing workflow.

Neither is the default. Agree on the first user task or workflow trigger and the evidence that proves delivery. Producing one video is an integration check.

Module 1. Choices and trade-offs

Match capability to risk and value

Choose the media format after the application goal:

- Batch avatar video: a rendering step in a content-production application; requires job handling and controlled publication.
- Real-time avatar: supports interaction; plan for latency, consent, moderation, and live support.
- Voice or audio-first: needs less media production; compare it against accessibility and channel requirements.
- Video translation: maintain language versions in the content workflow; review translated meaning before publication.
- No avatar: use text or audio when a visible presenter adds no value or cannot meet the requirements.

The decision must cover consent, likeness and voice rights, supported regions, identity model, pricing, accessibility coverage, content-safety controls, and an exit path.

Module 1. What must be true

A capability decision the sponsor can defend

Record these decisions before provisioning:

- selected experience capability and why it fits the pilot
- application goal, user task or workflow trigger, and acceptance checks
- alternatives considered and why they were not selected
- consent, disclosure, accessibility, privacy, residency, and retention assumptions
- operating owner, support path, and conditions that would stop or simplify the experience

Discussion: If the avatar option became unavailable tomorrow, could the onboarding outcome still be delivered safely?

Module 2. Context

Provision the Foundry and Speech foundation

Connect the selected services through the customer's approved identity and network.

The foundation should support:

- keyless access patterns where possible
- a model and grounding path only when assisted authoring or live answers are needed
- Speech or media services for the selected experience capability
- access to the approved content; add indexed retrieval only when the selected path needs it
- telemetry that helps owners understand behavior without over-collecting employee data

Module 2. Choices and trade-offs

Keep components replaceable

Agree these platform choices:

- Identity: managed identity and role-based access reduce secret handling, but require clear ownership.
- Region and capacity: availability, latency, data residency, and cost may point to different deployment choices.
- Model path: a general model may be enough for drafting; stricter use cases may need more evaluation and guardrails.
- Observability: useful traces help diagnose issues, but message content capture must be intentional and governed.
- Integration contracts: keep rendering, assistant, content, and approval components replaceable.

Module 2. What must be true

A foundation ready for governed work

Verify these conditions:

- required resources are provisioned with the intended identity model
- environment settings are documented without exposing secrets
- content, assistant, rendering, approval, and telemetry components can connect
- owners know where logs, traces, and configuration evidence will live
- the runtime identity can reach the selected services from the customer's approved network

Discussion: Can the team explain who can access what, why, and how that access is reviewed?

Module 3. Context

Build the governed content pipeline

The avatar cannot become an unreviewed policy source. It must express approved content.

For every claim used in the experience, the pipeline needs:

- source, version, owner, and review cycle
- audience and locale
- sensitivity and escalation route
- expiry or withdrawal conditions
- traceability from source to script segment

This makes onboarding content a controlled input for the assistant, storyboard, approvals, and published experience.

Module 3. Choices and trade-offs

Govern claims without slowing every edit

Discussion choices:

- What counts as an approved source: policy page, HR guide, learning page, legal FAQ, or SME note?
- Which changes need full reapproval versus owner review?
- How granular should claims be for traceability?
- How are locale-specific policy differences handled?
- Who can retire, pause, or replace content when guidance changes?

Start with a small pilot corpus and name its content owner.

Module 3. What must be true

Versioned claims with named ownership

Build a traceable claim set that downstream steps can use:

- each claim has an authoritative source and owner
- versions and expiry rules are visible
- sensitive topics are marked with review and escalation requirements
- unapproved or expired material is excluded from drafting
- the source-to-script relationship can be shown to reviewers
- exact wording is approved before module 5 creates a private preview

Discussion: If an employee challenges a statement in the avatar experience, can the owner show where it came from and whether it was current?

Module 4. Context

Build the grounded assistant behind the experience

Skip model-based drafting when the wording is already approved. When assisted authoring or live answers are needed, connect only the permitted sources.

Expected behavior:

- cite approved sources when making claims
- refuse or escalate when content is missing or expired
- avoid policy interpretation beyond the approved corpus
- support human review rather than replacing it
- preserve a clear path from source to script to final experience

Reviewers must be able to see why the assistant said what it said.

Module 4. Choices and trade-offs

Helpful drafting versus unsafe authority

The assistant can speed preparation. It cannot become the decision-maker.

Design choices include:

- how strict retrieval should be before drafting
- whether the assistant answers employee questions directly or routes to human support
- how it handles missing, conflicting, or stale content
- what tone and reading level it uses for onboarding audiences
- how citations appear for reviewers versus employees

Test grounded drafting, refusal, and escalation before adding more supported questions.

Module 4. What must be true

Cited drafts and visible refusals

Check the generated drafts and live responses:

- draft segments include source links or claim references
- unsupported requests are refused or routed for human help
- reviewer prompts return source evidence for each claim
- sensitive policy questions trigger the expected escalation path
- logs or traces help diagnose grounding failures without overexposing employee data

Discussion: Which answer would worry us more: "I don't know" or a confident answer without a source?

Module 5. Context

Connect generation to the application

Content production needs a deployed workflow with persisted job state and failure recovery. An interactive assistant needs a client connected to the bounded answer path.

Check the experience with its intended users:

- synthetic-media disclosure is visible and plain
- captions and transcript are available
- keyboard, mobile, and low-bandwidth access are considered
- each language version receives a review against its approved source
- a non-avatar alternative exists where needed

Include accessibility and disclosure in the client and publishing workflow.

Module 5. Choices and trade-offs

Review presentation and access

Agree these experience choices:

- Avatar style: check realistic avatars for impersonation risks.
- Voice: confirm consent and usage rights before selecting a branded voice.
- Disclosure placement: early and visible disclosure is safer than hidden footnotes.
- Fallback: transcript, audio, or human-led alternatives reduce exclusion.
- Localization: translation must preserve meaning, policy nuance, and accessibility.
- Channel: Teams, learning platform, intranet, or email each changes measurement and support.

The experience should never imply that a real person said something they did not approve.

Module 5. What must be true

Approved revision rendered accessibly

Build an experience package that can be reviewed before release:

- generated only from an approved script revision
- includes disclosure, captions, transcript, and fallback
- preserves source and approval references
- supports the selected audience, locale, and channel
- identifies who can pause or withdraw the published version

For content production, repeat a trigger and recover an interrupted job without duplicate output. For an assistant, test supported and refused questions through the actual client.

Discussion: Could an employee understand that the media is synthetic, get the same message without the avatar, and find human help?

Module 6. Context

Gate publication behind human approval

Bind human approval to the exact revision that will reach users.

Before anything is published, named reviewers must check:

- factual accuracy and source alignment
- legal, compliance, privacy, and labor considerations
- brand, tone, and employee experience fit
- accessibility and inclusive design expectations
- publication scope, support route, and withdrawal authority

For content production, module 3 approved the wording; this module approves the media and release. For an assistant, approve the application version and content/refusal policy. Enforce the decision in the customer's release system.

Module 6. Choices and trade-offs

Keep approval strong and workable

Decisions to make:

- Which roles are mandatory for each content type?
- What changes reset approval: source change, translation, visual edit, voice change, or channel move?
- Who can emergency-pause an experience?
- How are reviewer disagreements resolved?
- What evidence must be retained, and for how long?

Verify that changed or unapproved revisions cannot publish.

Module 6. What must be true

Publication and withdrawal record

Create a release record that shows:

- approved script revision and rendered experience
- reviewer roles, decisions, and conditions
- accessibility and disclosure review outcome
- publication channel, audience, and owner
- pause, withdrawal, and replacement process

Prove withdrawal through the user channel. Changing a local approval file is insufficient.

Discussion: If a source policy changes after launch, who knows, who acts, and what happens to the published experience?

Module 7. Context

Evaluate, red-team, trace, and operate

Use the pilot's quality and user-task evidence to make the release decision.

Useful evidence includes:

- grounding quality and unsupported-claim defects
- disclosure and accessibility checks
- comprehension, task completion, and support handoffs
- feedback themes from employees and reviewers
- red-team findings around deception, bias, unsafe advice, and stale content
- trace review for failures and improvement opportunities

Use the evidence to decide whether to iterate, scale, pause, or withdraw.

Module 7. Choices and trade-offs

Use aggregate measurement

Agree the measurement rules:

- cohort-level insight is usually safer than individual-level monitoring
- qualitative feedback explains confusion that metrics can hide
- red-team scenarios should include accessibility, language, disclosure, and policy edge cases
- trace capture helps debugging but must respect privacy and data minimization
- scorecards should define action thresholds before the pilot starts

Keep onboarding measurement aggregate and identifier-free.

Module 7. What must be true

Application acceptance and release decision

Prove the deployed application. For content production, run a source update through generation and approved publication, then recover a failed job and withdraw outdated output. For an assistant, prove the supported task and refusal through its authenticated client.

Keep release evidence in the customer's existing system:

- pilot scorecard tied to the original onboarding outcome
- known defects and remediation owners
- red-team results and accepted residual risks
- trace review themes and operating improvements
- recommendation to iterate, scale, pause, or withdraw

Discussion: What evidence would convince this customer to expand the pattern to another topic, locale, or audience?

Close the discussion

Turn the decisions into a pilot plan

Bring one real onboarding moment and the people who own it to the next module.

Before the next working session, record:

- pilot topic, audience, locale, and channel
- application goal and first user task or workflow trigger
- approved source owners and review expectations
- selected experience capability and fallback
- disclosure, accessibility, and human-help requirements
- approval gate and withdrawal path
- first scorecard for operating evidence
- one owner and due date for every unresolved decision

The Agentic Co-build covers the bounded pilot. Production integration remains the customer's work.